

fc_net

February 2, 2023

```
[ ]: # %load fc_net.py
import numpy as np

from .layers import *
from .layer_utils import *

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity and
    softmax loss that uses a modular layer design. We assume an input dimension
    of D, a hidden dimension of H, and perform classification over C classes.

    The architecture should be affine - relu - affine - softmax.

    Note that this class does not implement gradient descent; instead, it
    will interact with a separate Solver object that is responsible for running
    optimization.

    The learnable parameters of the model are stored in the dictionary
    self.params that maps parameter names to numpy arrays.
    """

    def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
                  dropout=0, weight_scale=1e-3, reg=0.0):
        """
        Initialize a new network.

        Inputs:
        - input_dim: An integer giving the size of the input
        - hidden_dims: An integer giving the size of the hidden layer
        - num_classes: An integer giving the number of classes to classify
        - dropout: Scalar between 0 and 1 giving dropout strength.
        - weight_scale: Scalar giving the standard deviation for random
          initialization of the weights.
        - reg: Scalar giving L2 regularization strength.
        """
```

```

self.params = {}
self.reg = reg

# ===== #
# YOUR CODE HERE:
#   Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
#   self.params['W2'], self.params['b1'] and self.params['b2']. The
#   biases are initialized to zero and the weights are initialized
#   so that each parameter has mean 0 and standard deviation weight_scale.
#   The dimensions of W1 should be (input_dim, hidden_dim) and the
#   dimensions of W2 should be (hidden_dims, num_classes)
# ===== #
self.params['W1'] = np.random.randn(input_dim, hidden_dims) * weight_scale
self.params['W2'] = np.random.randn(hidden_dims, num_classes) * weight_scale
self.params['b1'] = np.zeros((1,hidden_dims))
self.params['b2'] = np.zeros((1,num_classes))
pass

# ===== #
# END YOUR CODE HERE
# ===== #

def loss(self, X, y=None):
    """
    Compute loss and gradient for a minibatch of data.

    Inputs:
    - X: Array of input data of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,). y[i] gives the label for X[i].

    Returns:
    If y is None, then run a test-time forward pass of the model and return:
    - scores: Array of shape (N, C) giving classification scores, where
      scores[i, c] is the classification score for X[i] and class c.

    If y is not None, then run a training-time forward and backward pass and
    return a tuple of:
    - loss: Scalar value giving the loss
    - grads: Dictionary with the same keys as self.params, mapping parameter
      names to gradients of the loss with respect to those parameters.
    """
    scores = None

    # ===== #
    # YOUR CODE HERE:
    #   Implement the forward pass of the two-layer neural network. Store
    #   the class scores as the variable 'scores'. Be sure to use the layers

```

```

# you prior implemented.
# ===== #
W1, W2 = self.params['W1'], self.params['W2']
b1, b2 = self.params['b1'], self.params['b2']
output_first_layer, cache_first = affine_relu_forward(X, W1, b1)
scores, cache_sec = affine_forward(output_first_layer, W2, b2)
pass
# ===== #
# END YOUR CODE HERE
# ===== #

# If y is None then we are in test mode so just return scores
if y is None:
    return scores

loss, grads = 0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backward pass of the two-layer neural net. Store
# the loss as the variable 'loss' and store the gradients in the
# 'grads' dictionary. For the grads dictionary, grads['W1'] holds
# the gradient for W1, grads['b1'] holds the gradient for b1, etc.
# i.e., grads[k] holds the gradient for self.params[k].
#
# Add L2 regularization, where there is an added cost  $0.5 * \text{self.reg} * W^2$ 
# for each W. Be sure to include the 0.5 multiplying factor to
# match our implementation.
#
# And be sure to use the layers you prior implemented.
# ===== #
loss, dout = softmax_loss(scores, y)
d_first_hidden, dW2, db2 = affine_backward(dout, cache_sec)
dx, dW1, db1 = affine_relu_backward(d_first_hidden, cache_first)
dW1 = dW1 + self.reg * W1
dW2 = dW2 + self.reg * W2
grads['W1'], grads['W2'] = dW1, dW2
grads['b1'], grads['b2'] = db1, db2
loss = loss + 0.5 * np.sum( self.reg * W1 * W1 ) + 0.5 * np.sum( self.reg *
↪ W2 * W2 )

pass

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```

```

class FullyConnectedNet(object):
    """
    A fully-connected neural network with an arbitrary number of hidden layers,
    ReLU nonlinearities, and a softmax loss function. This will also implement
    dropout and batch normalization as options. For a network with L layers,
    the architecture will be

    {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax

    where batch normalization and dropout are optional, and the {...} block is
    repeated L - 1 times.

    Similar to the TwoLayerNet above, learnable parameters are stored in the
    self.params dictionary and will be learned using the Solver class.
    """

    def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
                  dropout=0, use_batchnorm=False, reg=0.0,
                  weight_scale=1e-2, dtype=np.float32, seed=None):
        """
        Initialize a new FullyConnectedNet.

        Inputs:
        - hidden_dims: A list of integers giving the size of each hidden layer.
        - input_dim: An integer giving the size of the input.
        - num_classes: An integer giving the number of classes to classify.
        - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
          the network should not use dropout at all.
        - use_batchnorm: Whether or not the network should use batch normalization.
        - reg: Scalar giving L2 regularization strength.
        - weight_scale: Scalar giving the standard deviation for random
          initialization of the weights.
        - dtype: A numpy datatype object; all computations will be performed using
          this datatype. float32 is faster but less accurate, so you should use
          float64 for numeric gradient checking.
        - seed: If not None, then pass this random seed to the dropout layers. This
          will make the dropout layers deterministic so we can gradient check the
          model.
        """
        self.use_batchnorm = use_batchnorm
        self.use_dropout = dropout > 0
        self.reg = reg
        self.num_layers = 1 + len(hidden_dims)
        self.dtype = dtype
        self.params = {}

```

```

# ===== #
# YOUR CODE HERE:
# Initialize all parameters of the network in the self.params dictionary.
# The weights and biases of layer 1 are W1 and b1; and in general the
# weights and biases of layer i are Wi and bi. The
# biases are initialized to zero and the weights are initialized
# so that each parameter has mean 0 and standard deviation weight_scale.
# ===== #
self.params['W1'] = np.random.randn(input_dim, hidden_dims[0]) * ␣
↪weight_scale
self.params['b1'] = np.zeros(hidden_dims[0])

for i in range(self.num_layers - 2):
    j = i + 2
    weights = str('W') + str(j)
    bias = str('b') + str(j)
    # print(weights, bias)
    self.params[weights] = np.random.randn(hidden_dims[j-2], ␣
↪hidden_dims[j-1]) * weight_scale
    self.params[bias] = np.zeros((1, hidden_dims[j-1]))

weight_last = str('W') + str(self.num_layers)
bias_last = str('b') + str(self.num_layers)
self.params[weight_last] = np.random.randn(hidden_dims[-1], num_classes)
self.params[bias_last] = np.zeros(num_classes)
pass

# ===== #
# END YOUR CODE HERE
# ===== #

# When using dropout we need to pass a dropout_param dictionary to each
# dropout layer so that the layer knows the dropout probability and the mode
# (train / test). You can pass the same dropout_param to each dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
    if seed is not None:
        self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running means and
# variances, so we need to pass a special bn_param object to each batch
# normalization layer. You should pass self.bn_params[0] to the forward pass
# of the first batch normalization layer, self.bn_params[1] to the forward

```

```

    # pass of the second batch normalization layer, etc.
    self.bn_params = []
    if self.use_batchnorm:
        self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]

    # Cast all parameters to the correct datatype
    for k, v in self.params.items():
        self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param since they
    # behave differently during training and testing.
    if self.dropout_param is not None:
        self.dropout_param['mode'] = mode
    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param[mode] = mode

    scores = None
    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the FC net and store the output
    # scores as the variable "scores".
    # ===== #
    output = {}
    output[0] = X
    cache = {}
    for i in range(self.num_layers):
        j = i + 1
        weights = 'W' + str(j)
        bias = 'b' + str(j)
        W_j = self.params[weights]
        b_j = self.params[bias]

        if j == self.num_layers:
            output[j], cache[i] = affine_forward(output[i], W_j, b_j)
        else:

```

```

        output[j], cache[i] = affine_relu_forward(output[i], W_j, b_j)

scores = output[self.num_layers]
pass

# ===== #
# END YOUR CODE HERE
# ===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backwards pass of the FC net and store the gradients
# in the grads dict, so that grads[k] is the gradient of self.params[k]
# Be sure your L2 regularization includes a 0.5 factor.
# ===== #
loss, dx = softmax_loss(scores, y)
reg_loss = 0
for i in range(self.num_layers):
    j = i + 1
    weights = str('W') + str(j)
    W_j = self.params[weights]
    reg_loss = reg_loss + 0.5 * self.reg * np.sum(W_j * W_j)
loss = loss + reg_loss

# compute grads
doutput = {}
hidden_dims = self.num_layers - 1
weight_last = str('W') + str(self.num_layers)
bias_last = str('b') + str(self.num_layers)
W_last, b_last = self.params[weight_last], self.params[bias_last]
doutput[hidden_dims], grads[weight_last], grads[bias_last] = \
    affine_backward(dx, cache[hidden_dims])
grads[weight_last] = grads[weight_last] + self.reg * self.
params[weight_last]

for i in range(hidden_dims):
    weight = str('W') + str(hidden_dims - i)
    bias = str('b') + str(hidden_dims - i)
    doutput[hidden_dims - i - 1], grads[weight], grads[bias] = \
    affine_relu_backward(doutput[hidden_dims - i], cache[hidden_dims - i - 1])
    grads[weight] = grads[weight] + self.reg * self.params[weight]

```

```
pass
```

```
# ===== #
```

```
# END YOUR CODE HERE
```

```
# ===== #
```

```
return loss, grads
```