

# neural\_net

February 2, 2023

```
[ ]: # %load neural_net.py
import numpy as np
import matplotlib.pyplot as plt

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network. The net has an input dimension of
    D, a hidden layer dimension of H, and performs classification over C classes.
    We train the network with a softmax loss function and L2 regularization on the
    weight matrices. The network uses a ReLU nonlinearity after the first fully
    connected layer.

    In other words, the network has the following architecture:

    input - fully connected layer - ReLU - fully connected layer - softmax

    The outputs of the second fully-connected layer are the scores for each class.
    """

    def __init__(self, input_size, hidden_size, output_size, std=1e-4):
        """
        Initialize the model. Weights are initialized to small random values and
        biases are initialized to zero. Weights and biases are stored in the
        variable self.params, which is a dictionary with the following keys:

        W1: First layer weights; has shape (H, D)
        b1: First layer biases; has shape (H,)
        W2: Second layer weights; has shape (C, H)
        b2: Second layer biases; has shape (C,)

        Inputs:
        - input_size: The dimension D of the input data.
        - hidden_size: The number of neurons H in the hidden layer.
        - output_size: The number of classes C.
        """
        self.params = {}
```

```

self.params['W1'] = std * np.random.randn(hidden_size, input_size)
self.params['b1'] = np.zeros(hidden_size)
self.params['W2'] = std * np.random.randn(output_size, hidden_size)
self.params['b2'] = np.zeros(output_size)

def loss(self, X, y=None, reg=0.0):
    """
    Compute the loss and gradients for a two layer fully connected neural
    network.

    Inputs:
    - X: Input data of shape (N, D). Each X[i] is a training sample.
    - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
        an integer in the range 0 <= y[i] < C. This parameter is optional; if it
        is not passed then we only return scores, and if it is passed then we
        instead return the loss and gradients.
    - reg: Regularization strength.

    Returns:
    If y is None, return a matrix scores of shape (N, C) where scores[i, c] is
    the score for class c on input X[i].

    If y is not None, instead return a tuple of:
    - loss: Loss (data loss and regularization loss) for this batch of training
        samples.
    - grads: Dictionary mapping parameter names to gradients of those parameters
        with respect to the loss function; has the same keys as self.params.
    """
    # Unpack variables from the params dictionary
    W1, b1 = self.params['W1'], self.params['b1']
    W2, b2 = self.params['W2'], self.params['b2']
    N, D = X.shape

    # Compute the forward pass
    scores = None

    # ===== #
    # YOUR CODE HERE:
    # Calculate the output scores of the neural network. The result
    # should be (N, C). As stated in the description for this class,
    # there should not be a ReLU layer after the second FC layer.
    # The output of the second FC layer is the output scores. Do not
    # use a for loop in your implementation.
    # ===== #
    X_after_first_layer = np.dot(X, W1.T) + b1 #broadcast
    X_after_ReLu = np.maximum(0, X_after_first_layer)

```

```

scores = np.dot(X_after_ReLu, W2.T) + b2
pass

# ===== #
# END YOUR CODE HERE
# ===== #

# If the targets are not given then jump out, we're done
if y is None:
    return scores

# Compute the loss
loss = None

# ===== #
# YOUR CODE HERE:
#   Calculate the loss of the neural network. This includes the
#   softmax loss and the L2 regularization for W1 and W2. Store the
#   total loss in teh variable loss. Multiply the regularization
#   loss by 0.5 (in addition to the factor reg).
# ===== #
N,C = scores.shape[0],scores.shape[1]
H,D = self.params['W1'].shape
scores = scores - np.max(scores, axis = 1, keepdims = True)

true_matrix = np.zeros((N,C))
true_matrix[np.arange(N), y ] = 1
true_score_mat = np.multiply(true_matrix, scores)
true_score = np.sum(true_score_mat)

false_score = np.log( np.sum( np.exp(scores), axis = 1 ) ).sum()

loss = false_score - true_score
loss = loss / N + 0.5 * reg * np.sum(W1**2) + 0.5 * reg * np.sum(W2**2)
# scores is num_examples by num_classes
pass
# ===== #
# END YOUR CODE HERE
# ===== #

grads = {}

# ===== #
# YOUR CODE HERE:
#   Implement the backward pass. Compute the derivatives of the
#   weights and the biases. Store the results in the grads

```

```

# dictionary. e.g., grads['W1'] should store the gradient for
# W1, and be of the same size as W1.
# ===== #
d_scores = np.zeros((N,C))
d_scores = np.exp(scores) / np.exp(scores).sum(axis = 1, keepdims = True)
d_scores = (d_scores - true_matrix) / N
grads['W2'] = ( X_after_ReLu.T.dot(d_scores) + reg * W2.T ).T
grads['b2'] = np.sum(d_scores, axis = 0).T

d_input2 = np.zeros((H,D))
d_input2 = np.dot(d_scores, W2)
d_input2[X_after_ReLu < 1e-5] = 0
# print(d_input2.shape)
grads['b1'] = np.sum(d_input2,axis = 0).T
grads['W1'] = (X.T.dot(d_input2) + reg * W1.T).T
# print(grads['W1'].shape, grads['b1'].shape, grads['W2'].shape,
↪grads['b2'].shape)

pass

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

def train(self, X, y, X_val, y_val,
        learning_rate=1e-3, learning_rate_decay=0.95,
        reg=1e-5, num_iters=100,
        batch_size=200, verbose=False):
    """
    Train this neural network using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) giving training data.
    - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
        X[i] has label c, where 0 ≤ c < C.
    - X_val: A numpy array of shape (N_val, D) giving validation data.
    - y_val: A numpy array of shape (N_val,) giving validation labels.
    - learning_rate: Scalar giving learning rate for optimization.
    - learning_rate_decay: Scalar giving factor used to decay the learning rate
        after each epoch.
    - reg: Scalar giving regularization strength.
    - num_iters: Number of steps to take when optimizing.
    - batch_size: Number of training examples to use per step.
    - verbose: boolean; if true print progress during optimization.
    """

```

```

num_train = X.shape[0]
num_val = X_val.shape[0]
iterations_per_epoch = max(num_train / batch_size, 1)

# Use SGD to optimize the parameters in self.model
loss_history = []
train_acc_history = []
val_acc_history = []

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

    # ===== #
    # YOUR CODE HERE:
    # Create a minibatch by sampling batch_size samples randomly.
    # ===== #
    idx = np.random.choice(range(num_train), batch_size)
    X_batch = X[idx]
    y_batch = y[idx]
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # Compute loss and gradients using the current minibatch
    loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
    loss_history.append(loss)

    # ===== #
    # YOUR CODE HERE:
    # Perform a gradient descent step using the minibatch to update
    # all parameters (i.e., W1, W2, b1, and b2).
    # ===== #
    self.params['W1'] = self.params['W1'] - learning_rate * grads['W1']
    self.params['W2'] = self.params['W2'] - learning_rate * grads['W2']
    self.params['b1'] = self.params['b1'] - learning_rate * grads['b1']
    self.params['b2'] = self.params['b2'] - learning_rate * grads['b2']
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    if verbose and it % 100 == 0:
        print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

    # Every epoch, check train and val accuracy and decay learning rate.
    if it % iterations_per_epoch == 0:
        # Check accuracy

```

```

train_acc = (self.predict(X_batch) == y_batch).mean()
val_acc = (self.predict(X_val) == y_val).mean()
train_acc_history.append(train_acc)
val_acc_history.append(val_acc)

# Decay learning rate
learning_rate *= learning_rate_decay

return {
    'loss_history': loss_history,
    'train_acc_history': train_acc_history,
    'val_acc_history': val_acc_history,
}

def predict(self, X):
    """
    Use the trained weights of this two-layer network to predict labels for
    data points. For each data point we predict scores for each of the  $C$ 
    classes, and assign each data point to the class with the highest score.

    Inputs:
    -  $X$ : A numpy array of shape  $(N, D)$  giving  $N$   $D$ -dimensional data points to
      classify.

    Returns:
    -  $y_{\text{pred}}$ : A numpy array of shape  $(N,)$  giving predicted labels for each of
      the elements of  $X$ . For all  $i$ ,  $y_{\text{pred}}[i] = c$  means that  $X[i]$  is predicted
      to have class  $c$ , where  $0 \leq c < C$ .
    """
    y_pred = None
    # ===== #
    # YOUR CODE HERE:
    # Predict the class given the input data.
    # ===== #
    X_after_first_layer = np.dot(X, self.params['W1'].T) + self.params['b1']
    ↪ #broadcast
    X_after_ReLu = np.maximum(0, X_after_first_layer)
    scores = np.dot(X_after_ReLu, self.params['W2'].T) + self.params['b2']
    y_pred = np.argmax(scores, axis = 1)
    pass

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred

```