

fc_net

February 10, 2023

```
[ ]: # %load fc_net.py
import numpy as np
import pdb

from .layers import *
from .layer_utils import *

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity and
    softmax loss that uses a modular layer design. We assume an input dimension
    of  $D$ , a hidden dimension of  $H$ , and perform classification over  $C$  classes.

    The architecture should be affine - relu - affine - softmax.

    Note that this class does not implement gradient descent; instead, it
    will interact with a separate Solver object that is responsible for running
    optimization.

    The learnable parameters of the model are stored in the dictionary
    self.params that maps parameter names to numpy arrays.
    """

    def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
                  dropout=1, weight_scale=1e-3, reg=0.0):
        """
        Initialize a new network.

        Inputs:
        - input_dim: An integer giving the size of the input
        - hidden_dims: An integer giving the size of the hidden layer
        - num_classes: An integer giving the number of classes to classify
        - dropout: Scalar between 0 and 1 giving dropout strength.
        - weight_scale: Scalar giving the standard deviation for random
          initialization of the weights.
        - reg: Scalar giving L2 regularization strength.
        """
```

```

"""
self.params = {}
self.reg = reg

# ===== #
# YOUR CODE HERE:
# Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
# self.params['W2'], self.params['b1'] and self.params['b2']. The
# biases are initialized to zero and the weights are initialized
# so that each parameter has mean 0 and standard deviation 1/weight_scale.
# The dimensions of W1 should be (input_dim, hidden_dim) and the
# dimensions of W2 should be (hidden_dims, num_classes)
# ===== #
W1, W2 = self.params['W1'], self.params['W2']
b1, b2 = self.params['b1'], self.params['b2']
output_first_layer, cache_first = affine_relu_forward(X, W1, b1)
scores, cache_sec = affine_forward(output_first_layer, W2, b2)

# ===== #
# END YOUR CODE HERE
# ===== #

def loss(self, X, y=None):
    """
    Compute loss and gradient for a minibatch of data.

    Inputs:
    - X: Array of input data of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,). y[i] gives the label for X[i].

    Returns:
    If y is None, then run a test-time forward pass of the model and return:
    - scores: Array of shape (N, C) giving classification scores, where
      scores[i, c] is the classification score for X[i] and class c.

    If y is not None, then run a training-time forward and backward pass and
    return a tuple of:
    - loss: Scalar value giving the loss
    - grads: Dictionary with the same keys as self.params, mapping parameter
      names to gradients of the loss with respect to those parameters.
    """
    scores = None

    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the two-layer neural network. Store

```

```

# the class scores as the variable 'scores'. Be sure to use the
layers
# you prior implemented.
# ===== #

# ===== #
# END YOUR CODE HERE
# ===== #

# If y is None then we are in test mode so just return scores
if y is None:
    return scores

loss, grads = 0, {}
# ===== #
# YOUR CODE HERE:
# Implement the backward pass of the two-layer neural net. Store
# the loss as the variable 'loss' and store the gradients in the
# 'grads' dictionary. For the grads dictionary, grads['W1'] holds
# the gradient for W1, grads['b1'] holds the gradient for b1, etc.
# i.e., grads[k] holds the gradient for self.params[k].
#
# Add L2 regularization, where there is an added cost  $0.5 * \text{self.reg} * W^2$ 
# for each W. Be sure to include the 0.5 multiplying factor to
# match our implementation.
#
# And be sure to use the layers you prior implemented.
# ===== #

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grads

```

```

class FullyConnectedNet(object):

```

```

    """

```

A fully-connected neural network with an arbitrary number of hidden layers, ReLU nonlinearities, and a softmax loss function. This will also implement dropout and batch normalization as options. For a network with L layers, the architecture will be

{affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax

where batch normalization and dropout are optional, and the {...} block is repeated L - 1 times.

Similar to the `TwoLayerNet` above, learnable parameters are stored in the `self.params` dictionary and will be learned using the `Solver` class.

```

"""
def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
              dropout=1, use_batchnorm=False, reg=0.0,
              weight_scale=1e-2, dtype=np.float32, seed=None):
    """
    Initialize a new FullyConnectedNet.

    Inputs:
    - hidden_dims: A list of integers giving the size of each hidden layer.
    - input_dim: An integer giving the size of the input.
    - num_classes: An integer giving the number of classes to classify.
    - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=1
    then
        the network should not use dropout at all.
    - use_batchnorm: Whether or not the network should use batch
    normalization.
    - reg: Scalar giving L2 regularization strength.
    - weight_scale: Scalar giving the standard deviation for random
        initialization of the weights.
    - dtype: A numpy datatype object; all computations will be performed
    using
        this datatype. float32 is faster but less accurate, so you should use
        float64 for numeric gradient checking.
    - seed: If not None, then pass this random seed to the dropout layers.
    This
        will make the dropout layers deterministic so we can gradient check
    the
        model.
    """
    self.use_batchnorm = use_batchnorm
    self.use_dropout = dropout < 1
    self.reg = reg
    self.num_layers = 1 + len(hidden_dims)
    self.dtype = dtype
    self.params = {}

    # ===== #
    # YOUR CODE HERE:
    # Initialize all parameters of the network in the self.params
    dictionary.
    # The weights and biases of layer 1 are W1 and b1; and in general the
    # weights and biases of layer i are Wi and bi. The

```

```

# biases are initialized to zero and the weights are initialized
# so that each parameter has mean 0 and standard deviation
↪weight_scale.
#
# BATCHNORM: Initialize the gammas of each layer to 1 and the beta
# parameters to zero. The gamma and beta parameters for layer 1
↪should
# be self.params['gamma1'] and self.params['beta1']. For layer 2,
↪they
# should be gamma2 and beta2, etc. Only use batchnorm if self.
↪use_batchnorm
# is true and DO NOT do batch normalize the output scores.
# =====
hidd_len = len(hidden_dims)
shape_one = input_dim
for i in range(hidd_len):
    self.params['W'+str(i+1)] = np.random.randn(shape_one,
↪hidden_dims[i]) * weight_scale
    self.params['b'+str(i+1)] = np.zeros(hidden_dims[i])
    shape_one = hidden_dims[i]
    if self.use_batchnorm is True:
        self.params['gamma'+str(i+1)] = np.ones(hidden_dims[i])
        self.params['beta'+str(i+1)] = np.zeros(hidden_dims[i])
    self.params['W'+str(hidd_len+1)] = np.random.randn(shape_one,
↪num_classes) * weight_scale
    self.params['b'+str(hidd_len+1)] = np.zeros(num_classes)
pass

# ===== #
# END YOUR CODE HERE
# ===== #

# When using dropout we need to pass a dropout_param dictionary to each
# dropout layer so that the layer knows the dropout probability and the
↪mode
# (train / test). You can pass the same dropout_param to each dropout
↪layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
if seed is not None:
    self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running means and
# variances, so we need to pass a special bn_param object to each batch

```

```

        # normalization layer. You should pass self.bn_params[0] to the forward
    ↪pass
        # of the first batch normalization layer, self.bn_params[1] to the
    ↪forward
        # pass of the second batch normalization layer, etc.
        self.bn_params = []
        if self.use_batchnorm:
            self.bn_params = [{'mode': 'train'} for i in np.arange(self.
    ↪num_layers - 1)]

        # Cast all parameters to the correct datatype
        for k, v in self.params.items():
            self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param since they
    # behave differently during training and testing.
    if self.dropout_param is not None:
        self.dropout_param['mode'] = mode
    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param['mode'] = mode

    scores = None

    # ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the FC net and store the output
    # scores as the variable "scores".
    #
    # BATCHNORM: If self.use_batchnorm is true, insert a bathnorm layer
    # between the affine_forward and relu_forward layers. You may
    # also write an affine_batchnorm_relu() function in layer_utils.py.
    #
    # DROPOUT: If dropout is non-zero, insert a dropout layer after
    # every ReLU layer.
    # ===== #
    output = {}

```

```

output[0] = X
cache = {}
for i in range(self.num_layers):
    j = i + 1
    weights = str('W') + str(j)
    bias = str('b') + str(j)

    W_j = self.params[weights]
    b_j = self.params[bias]

    if j == self.num_layers:
        output[j], cache[i] = affine_forward(output[i], W_j, b_j)
    else:
        if self.use_dropout and self.use_batchnorm:
            gammas = str('gamma') + str(j)
            betas = str('beta') + str(j)
            gamma_j = self.params[gammas]
            beta_j = self.params[betas]
            output[j], cache[i] =
↪affine_batchnorm_relu_dropout_forward(output[i], W_j, b_j, gamma_j, beta_j,
↪self.bn_params[i], self.dropout_param)

            elif self.use_batchnorm is True:
                gammas = str('gamma') + str(j)
                betas = str('beta') + str(j)
                gamma_j = self.params[gammas]
                beta_j = self.params[betas]
                output[j], cache[i] = affine_batchnorm_relu_forward(output[i],
↪W_j, b_j, gamma_j, beta_j, self.bn_params[i])

            else:
                output[j], cache[i] = affine_relu_forward(output[i], W_j, b_j)

scores = output[self.num_layers]
pass

# ===== #
# END YOUR CODE HERE
# ===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
# ===== #
# YOUR CODE HERE:

```

```

# Implement the backwards pass of the FC net and store the gradients
# in the grads dict, so that grads[k] is the gradient of self.
↪params[k]
# Be sure your L2 regularization includes a 0.5 factor.
#
# BATCHNORM: Incorporate the backward pass of the batchnorm.
#
# DROPOUT: Incorporate the backward pass of dropout.
# ===== #

loss, dout = softmax_loss(scores, y)
reg_loss = 0
for i in range(self.num_layers):
    j = i + 1
    weights = str('W') + str(j)
    W_j = self.params[weights]
    reg_loss = reg_loss + 0.5 * self.reg * np.sum(W_j * W_j)
loss = loss + reg_loss

#compute grad
hidd_len = self.num_layers-1
# weights and bias of the last layer
dx, last_weight, last_bias = affine_backward(dout, cache[self.
↪num_layers-1])
grads['W'+str(self.num_layers)] = last_weight + self.reg * self.
↪params['W'+str(self.num_layers)]
grads['b'+str(self.num_layers)] = last_bias
# weights, bias, gamma, beta of the rest layers
if self.use_dropout and self.use_batchnorm:
    for i in range(hidd_len, 0, -1):
        dx, dw, db, dgamma, dbeta = ↪
↪affine_batchnorm_relu_dropout_backward(dx, cache[i-1])
        grads['W'+str(i)] = dw + self.reg * self.params['W'+str(i)]
        grads['b'+str(i)] = db
        grads['gamma'+str(i)] = dgamma
        grads['beta'+str(i)] = dbeta
elif self.use_batchnorm is True:
    for i in range(hidd_len, 0, -1):
        dx, dw, db, dgamma, dbeta = affine_batchnorm_relu_backward(dx, ↪
↪cache[i-1])
        grads['W'+str(i)] = dw + self.reg * self.params['W'+str(i)]
        grads['b'+str(i)] = db
        grads['gamma'+str(i)] = dgamma
        grads['beta'+str(i)] = dbeta
else:
    for i in range(self.num_layers-1, 0, -1):
        dx, dw, db = affine_relu_backward(dx, cache[i-1])

```



```
        grads['W'+str(i)] = dw + self.reg * self.params['W'+str(i)]
        grads['b'+str(i)] = db

    pass

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grads
```